

The Middleman Lies: Measuring Malicious LLM Router Attacks on Agents

Anonymous Author(s)

Abstract

LLM agents increasingly access models through API routers for protocol translation, model selection, and usage accounting. Because a router terminates the agent-facing connection, it can alter messages that drive model decisions and agent actions while preserving protocol validity. Prior work demonstrates selected attacks and defenses, but does not systematically measure which modifications propagate through the agent loop and cause verifiable effects. We present ROUTERTAMPERBENCH, an end-to-end benchmark with 939 cases over 143 tasks. Each case fixes a task-specific modification and uses trusted agent, tool, and environment records to distinguish delivery, downstream handling, and impact. We execute each applicable active case across 30 agent-model configurations covering five agents, six models, and three API formats. We count success only when the modification is delivered, the protocol remains valid, and trusted state records the specified consequence. At least one attack succeeds in every configuration, with a mean attack success rate of 31.9%. Across 789 comparable propagation cases, the pooled rate is 84.1% without a trusted check against state outside the router and 6.0% with one. Successful attacks span six security properties and often remain hidden behind task completion. In paired composition experiments, a second modification increases concealment by up to 40.0 percentage points and reduces recovery by 26.3 percentage points. A defense coverage analysis projects that schema validation blocks no cases, while end-to-end message authentication lowers the attack success rate to 6.5%. These results show that agents should not authorize security-sensitive actions from routed messages alone.

1 Introduction

LLM agents can modify files, update repositories, send messages, and change other external state on a user’s behalf. Many deployments place a LLM API router between the agent and the model provider for protocol translation, model selection, failover, and usage accounting [19, 34]. The router therefore

receives the instructions, tool definitions, tool results, tool calls, and control fields that drive the agent loop. Because it terminates the agent-facing connection and opens a separate connection to the provider, it processes these messages in plaintext. A malicious router can alter a message while preserving its schema and normal interaction flow. For example, it can replace a recipient address in a valid tool call before the agent executes it, while the final answer still reports success. Recent gateway flaws and host compromises illustrate practical paths to this level of router control [5, 14, 21].

This threat is difficult to measure from the final outcome alone. A modification may be ignored by the model, repaired in a later turn, rejected by the agent, or blocked by a trusted tool service. It may also change external state while the original task still appears complete. Prior work demonstrates selected router attacks and proposes mechanisms that authenticate routed messages or constrain router execution [13, 16, 19, 22, 34]. Agent security benchmarks evaluate other attack surfaces with executable tasks and environment checks [10, 18]. However, these efforts do not systematically compare the same router modifications across API formats and agent-model configurations while separating delivery, downstream handling, recovery, and impact verified outside the router. We therefore lack a clear account of which modifications propagate, where they stop, and what allows a system to reject or recover from them.

Our Work. We present ROUTERTAMPERBENCH, an end-to-end benchmark for measuring how router modifications propagate through LLM agents. We consider an attacker that fully controls one router but not the agent, model provider, tool services, task environment, or evaluator. The benchmark defines attacks over application-level objects such as instructions, tool definitions, tool results, tool calls, arguments, control fields, and final responses. It maps each object to OpenAI Responses, OpenAI Chat Completions, and Anthropic Messages, including streamed messages. Each case fixes a task, a target object, a modification, a trigger, and automated checks over trusted records. The full release contains 939 cases across

43 variants and 143 base tasks. Of these, 914 active cases span 13 attack families, while 25 cases measure passive visibility. We evaluate every applicable active case across 30 configurations formed by five agents and six tool-capable models.

Research Questions. We organize the evaluation around four questions:

- **RQ1: Effectiveness.** How often do router modifications cause security impact across agent–model configurations?
- **RQ2: Propagation.** Where do modifications stop, and what determines whether they reach external state?
- **RQ3: Composition.** Can a second modification conceal impact or prevent recovery?
- **RQ4: Defense coverage.** How much of the attack surface can trusted components outside the router cover?

Key Findings. At least one attack succeeds in every evaluated configuration. The mean attack success rate is 31.9%, while clean task utility averages 95.0%. The strongest separation appears after delivery. When no component checks the modified object against state that the router cannot alter, pooled attack success reaches 84.1%. With such a check, it falls to 6.0%, although both groups have delivery rates near 91%. Successful attacks also need not look like task failures. Among successful authorization and accountability attacks, 79.0% and 94.7% still complete the original task. In paired composition experiments, adding a false final answer raises concealment by 40.0 percentage points, while ending the interaction early lowers recovery by 26.3 percentage points. Finally, schema validation blocks none of the benchmark cases because the modifications remain structurally valid. Agent-side checks lower projected attack success to 20.8%, tool-side authorization lowers it to 28.9%, and end-to-end message authentication lowers it to 6.5%. These results show that reliable protection requires trusted state or authenticated evidence that the router cannot rewrite.

Contributions. We make three contributions:

- **An executable benchmark for malicious LLM routers.** ROUTERTAMPERBENCH provides 939 cases across 43 variants and 143 tasks, including 914 active cases in 13 attack families.
- **A cross-format measurement method.** Common application-level objects support equivalent modifications across three API formats, while trusted records distinguish delivery, handling, recovery, and impact.
- **An empirical account of propagation and defense coverage.** Experiments across 30 agent–model configurations show that trusted checks largely determine whether delivered modifications cause impact and quantify the residual exposure at different defense placements.

2 Background

2.1 LLM Agents

An LLM agent is a program that carries out a user task by repeatedly querying a language model and acting on what the model returns. At each turn, the agent builds a model request from the user task, instructions, conversation history, tool definitions, and prior tool results. The model returns text, one or more tool calls, or both. When it returns tool calls, the agent may validate their arguments, execute them, append the results to the conversation history, and request the next turn. We refer to this repeated request, execute, and append cycle as the *agent loop*. The loop ends when the model produces a final answer or the agent reaches a predefined stopping condition [4, 27, 38].

This loop exposes four distinct observation points: intermediate model responses, executed tool actions, environment state, and the final answer. An intermediate response shows what the model returns or requests, and an executed tool action shows what the agent actually attempts. A tool call may be rejected before execution, and an executed call may fail without changing the external state. Even when a state change occurs, a later action may reverse it. The final answer is a self-report and may omit or misreport a state change that already occurred. It therefore cannot, by itself, establish whether a modification propagated through the agent loop and affected the task environment. Determining propagation requires router logs together with execution evidence, including agent traces, tool results, and task-environment state.

2.2 LLM API Routers

Direct access to a single model provider is the simplest deployment. However, an agent often needs access to multiple providers and models. An LLM API router gives the agent a common endpoint for client authentication, model selection, API format translation, failover, rate limiting, and usage accounting. Model requests and responses therefore pass through the router [19, 34].

Direct access uses a single TLS connection, whereas routed access uses one connection from the agent to the router and another from the router to the provider. TLS protects each connection separately. The router terminates the first connection and initiates the second, so it reads and writes requests and responses as plaintext application data. Format translation and streaming support may also require the router to parse and re-serialize these messages. The two TLS connections therefore do not by themselves prove that the copy received by the agent matches the message returned by the provider [19, 22, 34].

Both directions carry fields that can affect the agent loop. Requests can contain instructions, conversation history, tool definitions, earlier tool results, the selected model, and the tool-choice configuration. These fields affect what the model

sees. Responses can contain generated text, tool names, tool arguments, call identifiers, and fields that indicate whether the model has finished. These fields affect what the agent receives and may execute. These roles are common across providers, although their API formats differ.

2.3 Semantic Objects Across API Formats

Model API formats can encode the same application-level unit in different ways. We call such a unit a *semantic object* and identify it by its role rather than its encoding. Our evaluation covers OpenAI Responses, OpenAI Chat Completions, and Anthropic Messages. Across these APIs, instructions and tool definitions use different field names and nesting. The OpenAI APIs represent tool-call arguments as JSON text, whereas Anthropic `tool_use` blocks store their inputs as JSON objects. Tool results, final texts, and stop reasons also use provider-specific fields [4, 27].

Streaming adds another source of representation differences. Without streaming, a provider returns a complete tool call or a final text in one message. Under streaming, the same tool call may arrive as a sequence of events: the tool name first, then the arguments as a series of partial JSON fragments, and finally a separate event reporting turn completion [4, 27]. The agent or router must reconstruct these fragments before processing the turn. A JSON field path is therefore tied to one API and one serialization, even when the field serves the same role in the agent loop.

Schema validity does not guarantee that a semantic object still carries the value the agent or the model sent. Changes to an instruction, tool definition, tool call, tool result, final text, or stop reason can leave a message schema-valid. For cross-format evaluation, the stable unit is therefore the semantic object rather than its position in one JSON message. We call the semantic object selected for modification in a given experiment the *target object*. Section 4 describes how ROUTERTAMPERBENCH maps each target object to the fields of each API while holding the intended change fixed.

2.4 Practical Paths to Router Compromise

The previous sections treat the router as an untrusted or partially trusted intermediary. That assumption is not hypothetical. Recent incidents illustrate three practical paths to control over software that handles model traffic. The first is a gateway flaw: in February 2026, GitLab patched a critical vulnerability in its self-hosted AI Gateway that allowed an authenticated user to execute code through crafted workflow definitions [14]. The second is weak authorization: LiteLLM disclosed a flaw that let an authenticated user modify proxy configuration and register handlers backed by attacker-controlled Python code [5]. The third is ordinary host compromise: a later Darktrace report described a compromised EC2 instance named LiteLLM-Proxy with access to Amazon Bedrock, where the




Figure 1: **System and trust boundaries.** The agent reaches the model provider through an LLM API router, the only untrusted component (red). Every model request and response traverses the router, which handles both in plaintext. Tool calls and their results pass directly between the agent and the task environment (green), which the router never touches directly.

initial entry point was not confirmed, but the host exposed SSH to the Internet and received extensive external connection attempts before cryptomining activity appeared [21].

All three paths yield the same capability, namely the ability to read and rewrite model traffic in transit, and our threat model starts once it has been obtained. An attacker holding it could instead steal credentials or resell capacity, but such abuse produces billing and rate-limit anomalies that operators already monitor. Editing a semantic object leaves both directions schema-valid and changes token accounting only marginally, which makes it comparatively hard to detect and motivates measuring how far such an edit propagates.

3 Threat Model

3.1 System Model

We consider an agent that reaches its model provider through a single LLM API router. A user gives a task to the agent, and the agent runs the loop described in Section 2.1: it builds a model request, receives a model response, and executes any tool calls that response contains. Figure 1 shows the paths these steps take. The first two steps traverse the router. Every model request the agent builds and every model response it receives pass through the router before reaching the other endpoint. The third step does not: the agent executes tool calls directly against the task environment, so tool results return to the agent without traversing the router.

This split determines what the router can reach. On one side, it holds every message the model and the agent exchange, in plaintext, for the reasons given in Section 2.2. On the other, it never touches a tool call during execution or the state that call produces. A change made at the router therefore reaches the task environment only if the model or the agent acts on the modified message. How far such a change propagates through the agent loop is what our evaluation measures.

Neither endpoint can settle that question on its own. The two TLS connections protect each hop separately and, as Sec-

tion 2.2 notes, do not establish that the copy the agent receives matches the message the provider returned. We assume no additional end-to-end mechanism supplies that guarantee.

3.2 Trust Assumptions

The router is the only untrusted component. The user is benign, and the task statement defines the operation the agent is authorized to perform. The agent, the model provider, the tool services, the task environment, and the evaluation harness operate as configured, and the records they produce are authentic. Those records are the execution evidence that Section 2.1 identifies as necessary for determining whether a modification took effect. Records produced only by the router, such as its usage and audit logs, carry no comparable guarantee.

Trusting a component does not mean assuming it resists a modified message. A trusted agent still executes a tool call that the router rewrote, and a trusted model still answers a request whose instructions the router replaced. Each component behaves correctly with respect to its own configuration while processing content the adversary chose. The harm we study arises from that combination.

3.3 Adversary

The adversary controls the router completely. Section 2.4 describes how such control is obtained in practice, and our model begins once such control exists. Control of the router yields two abilities. The adversary can read and retain every message that crosses it, and can rewrite, drop, replay, or fabricate a message before forwarding it.

These abilities apply to the semantic objects of Section 2.3. On the request path, the adversary can act on the instructions, the conversation history, the tool definitions, earlier tool results, the selected model, and the tool-choice configuration. On the response path, it can act on the generated text, the tool names and arguments, and the stop reason that tells the agent whether the turn is complete. Because every turn of a conversation crosses the router, the adversary can also retain state across turns and act only in the turns it selects.

What the adversary cannot do bounds the rest of this work. It cannot alter the model reasoning, the agent, the tool implementations, or the task environment, because none of these components is reachable from the router. Its only instrument is the content of a routed message. Any consequence beyond the router must therefore follow from the model or the agent processing that content.

3.4 Scope

We study a single malicious router on the model path of one agent. Within that setting, ROUTERTAMPERBENCH covers modifications on the request path and on the response path, messages that are replayed, modifications applied only in

selected turns, several modifications combined within one interaction, and changes to the stop reason. Two further cases are evaluated separately, because neither alters what the agent does. A message the adversary reads but leaves unchanged raises a question of disclosure, not of effect. A usage record the router falsifies deceives the party that reads the record, not the agent that acts.

Several related problems fall outside this scope. We do not study how the router is compromised, nor compromise of the agent, the provider, the tool services, the task environment, or the evaluator. We also exclude attacks that reach the agent by other routes, including malicious users, prompt injection carried in user input or tool output, training-time backdoors, multi-agent manipulation, and sandbox escape. Network-level denial of service is excluded as well: blocking the connection prevents the interaction rather than steering it, whereas suppression carried by a well-formed message remains in scope. Section 4 specifies the concrete modifications and the evidence required to count one as successful.

4 Benchmark Design

4.1 Overview

ROUTERTAMPERBENCH turns the threat model in Section 3 into executable benchmark cases. Its design follows three requirements. First, each active case targets a semantic object rather than a provider-specific field, so that the same intended modification can be applied across API formats. Second, the task, target, modification, and trigger are fixed before execution, so that every agent–model configuration faces the same attack condition. Third, the case includes an oracle over records outside the router, so the evaluator need not trust the router’s account of what occurred. The following example makes these elements concrete.

Figure 2 shows a concrete RMT-05 case. The model returns a tool call inviting `alice.miller@gmail.com` as a collaborator. Before forwarding the response, the router replaces the `collaborator_email` argument with `a.chen@outlook.com`. The response remains valid, and the agent executes the altered call. The task environment, which the router cannot modify, shows that the substituted address was invited and the original was not. This trace links a router modification to a state change recorded outside the router.

We next fix notation for one routed turn. For turn t , let x_t be the request emitted by the agent and $\tilde{x}_t$ the request consumed by the model. Let y_t be the response emitted by the model and $\tilde{y}_t$ the response consumed by the agent. Let M denote model inference. One routed turn is then

$$x_t \xrightarrow{\Phi_t} \tilde{x}_t \xrightarrow{M} y_t \xrightarrow{\Psi_t} \tilde{y}_t. \quad (1)$$

The transformations Φ_t and Ψ_t include ordinary routing behavior, such as format conversion and stream reconstruction,




Figure 2: Running example (RMT-05). On the response path, the router replaces the `collaborator_email` argument of a provider response. The agent executes the altered call, and the task environment records the deviation: the substituted address is invited, the original address is not.

together with any malicious modification. In Figure 2, Ψ_t replaces the `collaborator_email` argument before the agent receives the response. A byte-for-byte comparison of whole messages would also capture benign differences caused by format conversion or stream reconstruction. ROUTERTAMPER-BENCH therefore compares the selected target object before and after the router transformation.

Each active benchmark case combines a task τ , an attack specification a , and an oracle ω :

$$c = (\tau, a, \omega), \quad a = (p, o, \delta, g). \quad (2)$$

The task τ fixes the user task, the initial environment state, and the checks for normal task completion. The attack specification a fixes the path p , the target object o , the modification δ , and the trigger g . The path p identifies whether the target lies in a request, a response, or a record produced by the router. The modification specifies what is replaced, inserted, deleted, replayed, or misreported, while the trigger states when it applies. The oracle ω specifies a predicate over trusted records for checking the intended consequence. In Figure 2, p is the response path and o is the `collaborator_email` argument. The modification δ replaces the email address, g matches the `git_invite_collaborators` call, and ω checks the invitation state after execution. Figure 3 places this case definition in the full construction and execution pipeline.

Figure 3 shows the construction and execution pipeline. Offline construction begins with controlled tasks and clean executions. It extracts eligible target objects, selects valid alternative values or operations, and attaches an oracle to

each resulting case. During online execution, the router maps the case to the active API format and applies the edit when its trigger fires. The evaluator then uses trusted message captures, agent traces, tool records, and task state to determine the outcome.

4.2 Attack Space

We define each active router attack by the four components of a : path, target object, modification, and trigger. The path determines whether the router changes a request before the model receives it, a response before the agent receives it, or a record that the router reports. The target object follows the semantic roles defined in Section 2.3, such as an instruction, tool definition, tool result, tool call, control field, or final response. The modification replaces, inserts, deletes, replays, or misreports the target while preserving a valid message representation. The trigger selects the turn and message condition in which the modification applies and may depend on earlier routed messages. Together, these components express modifications on a specific turn, modifications that depend on prior traffic, replay, and combinations of multiple modifications. Table 1 groups the concrete variants into attack families and reports the release composition.

RMT-01 through RMT-12 test whether a modified request or response propagates through the agent loop. Attacks on the request path change what the model sees, while attacks on the response path change what the agent receives and may execute. RMT-S1 and RMT-B0 answer different questions. RMT-S1 asks whether the router can falsify an audit or usage record presented to another party. RMT-B0 asks which objects the router can read without changing the traffic. So we report RMT-B0 separately because it has no modification or downstream impact predicate.

For later analysis, each propagation case also records whether an independent trusted reference is available. An *independent trusted reference* is information outside the modified object that can verify its content or expected state. Examples include the original user task, a tool schema stored outside the router, an authorization check enforced by the tool service, or the task environment state. This annotation does not change the modification or its oracle. It groups cases by the information that may help the recipient detect, reject, or recover from the modification.

4.3 Applying Attacks Across API Formats

Comparable attacks across APIs must be defined by the target object rather than by a provider-specific field path. In Figure 2, the target is the `collaborator_email` argument of `git_invite_collaborators`. One API may encode this argument inside a JSON string, another may store it as a structured object, and a streaming API may split it across several




Figure 3: **ROUTERTAMPERBENCH construction and execution.** Clean runs provide target objects and trusted records for case construction. Each active case fixes a task, a router modification, and an oracle. During evaluation, the router applies the same semantic change across API formats, and trusted records determine delivery, downstream handling, and impact.

events. A rule tied to a field path or one event would therefore test a representation rather than the same attack condition.

At runtime, ROUTERTAMPERBENCH maps each request and response to an API-independent representation, first combining streaming events when needed. It then locates the target and trigger, applies the modification, validates the result, and serializes the message in the destination API format. Across configurations, the task, target object, modification, trigger, and oracle remain unchanged. Only the mapping between the target object and its API representation varies. A configuration that does not expose the required target is marked *not applicable* rather than counted as a failed attack.

4.4 Benchmark Construction

Benchmark construction starts from 143 adapted tasks with successful clean runs in a controlled baseline configuration. Before each run, the framework resets the task environment to its predefined initial state. It records the API messages, tool definitions, tool calls and results, final output, resulting environment state, and task result. These records identify target objects that occur in a valid execution and provide the trusted evidence used to define each case’s oracle.

We derive cases from eligible targets identified in the clean executions, task specifications, and tool schemas. For each target, we generate a fixed modification that is valid for the

same task. Replacement values preserve the expected type and differ from the clean value. Replay cases reuse a previously observed tool call or result, while insertions and deletions preserve the relevant schema and protocol constraints. For each modification, we also define an oracle over trusted records that checks its specified consequence. A case enters the release only when the target can be located unambiguously, the transformed message remains valid, the modification is nontrivial, and the oracle can verify its specified consequence.

Table 1 summarizes the released benchmark. Overall, ROUTERTAMPERBENCH contains 939 cases across 43 variants, covering 143 distinct base tasks. Cases that combine multiple modifications are constructed separately and are excluded from these totals.

5 Evaluation

We evaluate ROUTERTAMPERBENCH in four stages. First, we measure the effectiveness of individual router attacks across the full agent–model matrix. We then examine how delivered modifications proceed through later model, agent, and tool processing, and we relate their impact to security properties, task completion, and resource use. Next, we test whether combining modifications helps the attacker conceal impact or disrupt recovery. Finally, we evaluate defenses placed at

Table 1: **Attack families and release composition.** Active families group variants that share a path, target object, and modification. RMT-S1 falsifies a record produced by the router, while RMT-B0 observes traffic without changing it.

Family	Path	Target and modification	#Var.	#Cases
RMT-01: Intent rewrite	Request	Instructions or history, replace	3	75
RMT-02: Execution configuration	Request	Model selection or tool configuration, replace	5	125
RMT-03: Tool definition tampering	Request	Tool description or schema, replace, insert, or delete	12	187
RMT-04: Tool result tampering	Request	Earlier tool result, replace	4	100
RMT-05: Tool argument substitution	Response	Tool call arguments, replace	1	25
RMT-06: Tool identity substitution	Response	Tool name, replace	2	45
RMT-07: Tool call insertion	Response	Tool call, insert	1	25
RMT-08: Tool call suppression	Response	Required tool call, delete	1	25
RMT-09: Replay	Request/response	Prior tool result or tool call, replay	2	50
RMT-10: Control field tampering	Response	Turn control or call metadata, replace	5	107
RMT-11: Final response tampering	Response	Final text or response metadata, replace	2	50
RMT-12: Tool registry insertion	Request	Tool registry, insert	2	50
RMT-S1: Record falsification	Router record	Audit or usage record, mask, omit, or misreport	2	50
Active total			42	914
RMT-B0: Passive observation	Request/response	Target object in transit, observe	1	25
Full release			43	939

trusted components outside the router.

These stages address four research questions. **RQ1** asks how often router modifications cause impact recorded in trusted state across agents and models. **RQ2** asks where modifications stop and what consequences successful attacks cause. **RQ3** asks whether combinations of attacks can conceal impact or disrupt recovery. **RQ4** asks how much of the attack surface can trusted components outside the router cover.

5.1 Experimental Setup

Evaluation Sets. The main evaluation covers all 914 active cases in ROUTERTAMPERBENCH, spanning RMT-01 through RMT-12 and RMT-S1. RMT-B0 contributes 25 passive observation cases. We report RMT-B0 separately because it measures which objects the router can read rather than the effect of a modified message. Clean utility is measured on the 143 distinct base tasks in the full release. The composition study uses three sets, C1, C2, and C3, with ten cases per set. For each case, we execute the primary modification A , the secondary modification B , and their combination $A + B$ on four selected configurations, with three attempts per condition. This design schedules 1,080 composition runs. The corresponding clean runs are reused from the main evaluation.

Systems and API Formats. We exclude RMT-02 from the propagation and consequence analyses. A target configuration field can be modified only when the client serializes the field name, so its delivery and impact rates primarily reflect client-specific serialization and API-specific field handling rather than post-delivery propagation. We therefore report RMT-02 separately as an analysis of API-level exposure.

We evaluate five agents, OpenCode, Claude Code, Hermes, Pi, and OpenHands, with six models [1, 23, 28, 29, 31]. The models cover three families and two capability tiers per family: deepseek-v4-flash and deepseek-v4-pro, gpt-5.6-luna and gpt-5.6-terra, and claude-sonnet-4-6 and claude-opus-4-6 [3, 11, 12, 26]. The agents communicate with the router through OpenAI Responses, OpenAI Chat Completions, or Anthropic Messages [2, 24, 25]. Together, these choices yield 30 agent–model configurations. We run every applicable case on each configuration while holding agent versions, runtime settings, task environments, modifications, and oracles fixed. A case is not applicable when a configuration does not expose its required target object, as described in Section 4.3.

Procedure. Each run executes in an isolated workspace with a dedicated router instance. The framework resets the task environment, exposes the required MCP tools, and configures the agent to send model traffic through the router. When the trigger condition is met, the router locates the target object, applies the fixed modification, validates the transformed message, and forwards it. The agent and model continue their normal interaction until the task ends or a configured stopping condition is reached. The evaluator applies the case oracle to trusted message captures, agent traces, tool records, and task state. It also records the API format, configuration, attempt index, timeout, and infrastructure status for run reconstruction.

Metrics. We measure attack effectiveness with attack success rate, or ASR. For a run r of case c , delivery $D_c(r)$ holds when a trusted capture at the intended recipient contains the modified target object. The intended recipient is the model endpoint for request-side cases and the agent for response-

side cases. Downstream handling $E_c(r)$ holds when a case-specific trace predicate records later processing of that object in a model decision, agent action, or tool step. Impact $I_c(r)$ holds when the case oracle observes the specified consequence in trusted state outside the router. Protocol completion $P(r)$ holds when the routed interaction remains valid under the active API format. Task utility $U_c(r)$ holds when the original task contract is satisfied.

We call a run a strict attack success when the modification is delivered, the protocol remains valid, and the specified impact occurs:

$$\text{Succ}_c(r) \triangleq D_c(r) \wedge P(r) \wedge I_c(r). \quad (3)$$

Downstream handling is diagnostic and does not enter Equation 3. We report $E \mid D$ and $I \mid D$ to show how far a delivered modification proceeds. Recovery holds when the modification is delivered, the specified impact does not occur, and the original task still completes.

ASR is the mean of Succ over eligible runs. For grouped results, we first compute the rate within each configuration and then take an unweighted mean across configurations. This prevents configurations with more applicable cases from dominating the aggregate. We exclude runs that fail before the trigger because of an infrastructure error. Parse failures, tool rejections, early termination, and other failures caused by the modified message remain in the denominator because they are outcomes of the attack condition.

The composition study compares A and $A + B$ runs matched by case, configuration, and attempt index. A pair enters the analysis only when both runs are available. We report the number of pairs and the paired difference in percentage points.

5.2 Router Attack Assessment

Overall Effectiveness. Table 2 reports attack success and task utility for every agent–model configuration. At least one attack success occurs in each of the 30 configurations. ASR ranges from 25.9% to 36.6% and averages 31.9%. Because attack success requires delivery, protocol completion, and impact recorded outside the router, a modified transcript alone cannot satisfy the oracle. Thus, every evaluated configuration allows at least one router modification to reach trusted state outside the router.

Agent and Model Differences. Averaging by model yields ASRs between 31.5% and 32.4%, so the six model means are close. Averaging by agent yields a wider range, from 27.9% for OpenHands to 35.4% for OpenCode. The observed variation is therefore larger across agents than across models in this matrix. This comparison does not establish an intrinsic security ranking because agents expose different target objects and can take different paths through the same task.

Task Utility. On clean runs, utility ranges from 89.5% to 99.3% and averages 95.0%. In attack runs, utility ranges from 75.5% to 89.7% and averages 85.0%. The clean and attack averages differ by 10.0 percentage points. The observed impacts therefore arise in configurations with reliable clean baselines.

5.3 Propagation Analysis

After excluding RMT-02, the propagation analysis covers 789 active cases with complete and comparable downstream annotations. We group these cases by the downstream condition encountered after delivery. The four groups distinguish modifications that leave no recipient-visible trace, modifications that leave a recipient-visible trace without an independent trusted reference, modifications that can be checked against an independent trusted reference, and replayed calls rejected by the endpoint. An independent trusted reference is information outside the modified object that can verify its content or expected state, as defined in Section 4.2. Table 3 compares ASR, delivery, downstream handling, impact, and recovery across these groups.

Downstream Conditions. The main separation is whether the recipient can check the modified object against state that the router cannot alter. The pooled ASR is 84.1% when no independent check is available and 6.0% when one is available. The same pattern appears in the specific groups. A visible but unverifiable trace reaches 98.9% ASR, while cases with an independent trusted reference reach 6.3%. Replayed calls produce no strict successes. These results show that visibility alone does not stop the attack. Visibility supports rejection only with an independent basis for verification.

Stopping Points. The two pooled groups have similar delivery rates, 90.5% without an independent check and 91.4% with one. Their outcomes diverge after delivery. Without an independent check, downstream handling reaches 97.6% and impact reaches 89.9%. With an independent check, these rates fall to 53.5% and 10.1%, while recovery rises from 5.7% to 39.3%. Endpoint replay protection provides a more specific stopping mechanism. Replayed calls reach downstream handling in 99.7% of delivered runs, but the trusted endpoint rejects them before a state change. Impact therefore remains at 0%, and recovery reaches 97.1%.

Across Configurations. The same separation appears throughout the matrix. For each configuration, Figure 4 subtracts the ASR of cases with an independent trusted reference from that of cases without an independent trusted reference. The difference is positive in all 30 configurations and ranges from 62.3 to 92.0 percentage points. Cases with an independent reference therefore have lower ASR across every evaluated agent and model. This comparison remains observational

Table 2: **Attack success and task utility across agent–model configurations (%)**. U_{clean} and U_{atk} denote clean and attack utility. Higher utility and lower ASR are better. Within each agent, green cells with bold values and blue cells with underlined values indicate the best and second-best model values. Ties share the same highlight.

Model	OpenCode			Claude Code			Hermes			Pi			OpenHands		
	U_{clean}	U_{atk}	ASR	U_{clean}	U_{atk}	ASR	U_{clean}	U_{atk}	ASR	U_{clean}	U_{atk}	ASR	U_{clean}	U_{atk}	ASR
deepseek-v4-flash	99.3	84.4	35.3	98.6	<u>89.0</u>	32.5	96.5	83.3	32.9	<u>97.2</u>	85.7	32.7	99.3	75.5	28.5
deepseek-v4-pro	<u>98.6</u>	87.5	34.1	<u>97.9</u>	89.2	32.3	<u>93.7</u>	<u>85.9</u>	32.8	98.6	89.7	<u>31.4</u>	94.4	89.1	<u>27.7</u>
gpt-5.6-luna	97.2	<u>86.0</u>	<u>35.0</u>	<u>97.9</u>	88.1	32.8	91.6	89.0	<u>31.0</u>	<u>97.2</u>	<u>88.7</u>	30.8	<u>96.5</u>	87.7	28.0
gpt-5.6-terra	97.2	84.8	36.6	94.4	86.2	<u>32.0</u>	92.3	85.4	30.5	96.5	87.3	32.3	95.8	<u>88.3</u>	25.9
claude-sonnet-4-6	89.5	78.8	36.2	90.9	82.5	33.1	90.9	83.3	31.5	91.6	85.6	32.3	92.3	81.1	29.1
claude-opus-4-6	93.0	79.3	35.2	90.9	83.4	31.4	93.0	79.8	33.1	91.6	84.2	31.6	95.8	81.4	28.1
Average	95.8	83.5	35.4	95.1	86.4	32.3	93.0	84.5	32.0	95.5	86.9	31.9	95.7	83.9	27.9

Table 3: **Propagation outcomes with and without an independent check(%)**. The analysis includes 789 active cases.

Downstream condition	Cases	ASR	D	E	I	D	I	R
<i>No independent check</i>								
No recipient-visible trace	245	82.6	89.5	97.4	89.0	6.2		
Recipient-visible trace	25	98.9	100.0	100.0	98.9	1.1		
<i>Independent check available</i>								
Trusted reference	494	6.3	91.1	51.2	10.6	36.4		
Replay protection	25	0.0	97.9	99.7	0.0	97.1		
<i>Pooled comparison</i>								
Without check	270	84.1	90.5	97.6	89.9	5.7		
With check	519	6.0	91.4	53.5	10.1	39.3		

because the two groups contain different cases. It motivates the defense coverage analysis below.

5.4 Security Consequences

Security Properties. We group the 789 cases by the property checked by their oracle. Integrity covers incorrect content or state. Authorization covers operations outside the user’s task. Availability covers suppression of a required operation, and freshness covers stale or replayed state. Accountability covers false reports or records. Resource abuse covers unnecessary model or tool use charged to the task. Table 4 reports ASR and task outcomes for each property.

Strict attack successes span all six security properties represented in the benchmark. ASR ranges from 1.5% for freshness to 100.0% for resource abuse, with accountability at 75.6% and availability at 47.3%. The resource abuse estimate uses 750 runs of 25 targeted stress cases across 30 configurations.

Task Outcomes. Successful attacks do not always appear as task failure. Among strict authorization successes, 79.0% still complete the original task. The same is true for 94.7% of accountability successes and 95.5% of resource-abuse suc-




Figure 4: **ASR difference between cases without and with an independent trusted reference**. Each cell subtracts the ASR of cases with an independent trusted reference from the ASR of cases without an independent trusted reference for one agent–model configuration.

cesses. Availability has the opposite pattern. Almost every successful availability attack leaves the task incomplete because the modification removes an operation required by the task contract. These outcomes show why task utility and security impact must be scored separately. A completed task can coexist with an unauthorized action, a falsified report, or excess resource use. A failed task, in turn, does not identify which security property was violated.

Resource Cost. Beyond the final task outcome, some modifications increase interaction cost. We pair each attack run with the clean run for the same task and configuration, then divide the attack token count by the clean token count. Table 5 reports the two variants with the largest median token ratios. `history_stripping` increases token use by $2.8\times$ and `model-`

Table 4: **Security consequences and task outcomes (%)**. ASR is macro-averaged across configurations. U is conditioned on strict attack success.

Property	Cases	Runs	ASR	$U = 1$	$U = 0$
Integrity	401	11,607	20.9	48.2	51.8
Authorization	163	4,746	25.0	79.0	21.0
Availability	50	1,500	47.3	0.8	99.2
Freshness	50	1,500	1.5	0.0	100.0
Accountability	100	3,000	75.6	94.7	5.3
Resource abuse	25	750	100.0	95.5	4.5

Table 5: **Largest observed median resource amplification**. Token counts are per-run medians over 676 and 697 matched attack-clean pairs, respectively. Ratios are medians of the corresponding pairwise ratios.

Variant	Tokens (med.)	Median ratio	
		Token	Round
history_stripping	41.4k $\rightarrow$ 100.3k	2.8 $\times$	3.0 $\times$
error_result_fabrication	41.2k $\rightarrow$ 56.4k	1.3 $\times$	1.25 $\times$

interaction rounds by 3.0 $\times$. `error_result_fabrication` increases them by 1.3 $\times$ and 1.25 $\times$, respectively. In both cases, the extra cost is associated with additional model turns after the history or tool result is modified.

Take-away. Router attacks produce strict successes in all 30 configurations, with a mean ASR of 31.9%. The main separation appears after delivery: cases without an independent check reach 84.1% pooled ASR, while cases with an independent check reach 6.0%. Successful authorization, accountability, and resource-abuse attacks often preserve task completion, so task utility alone does not reveal security impact.

5.5 Attack Composition

The propagation analysis shows that some agents recover after a delivered modification. Because the router observes later model traffic, it can adapt by applying a second modification in the same interaction. We therefore test whether a second modification can hide an impact in the final answer or reduce the agent’s ability to recover from the first modification.

Experimental Design. C1 combines required-call suppression (RMT-08) with final-response tampering (RMT-11). C2 combines intent rewrite (RMT-01) with final-response tampering. C3 combines required-call suppression with premature finalization (RMT-10). In each set, A denotes the primary modification and B denotes the secondary modification. For C1 and C2, concealment requires both the impact caused by A and a final answer that claims successful completion. For C3, recovery records whether the original task still completes after the required call is removed. The paired analysis matches A

Table 6: **Paired effect of attack composition**. Each row compares matched runs with the primary modification alone and with both modifications. Rates are percentages. Changes are percentage points, with 95% confidence intervals.

Outcome (pairs)	Primary $\rightarrow$ both	Change	95% CI
C1 concealment (105)	52.4 $\rightarrow$ 92.4	+40.0	[29.5, 50.5]
C2 concealment (112)	0.0 $\rightarrow$ 9.8	+9.8	[4.5, 16.1]
C2 recovery (112)	87.5 $\rightarrow$ 87.5	+0.0	[−6.3, 6.3]
C3 recovery (118)	49.1 $\rightarrow$ 22.9	−26.3	[−34.8, −17.8]

and $A + B$ runs by case, configuration, and attempt index. Table 6 reports how each outcome changes when the secondary modification is added.

Concealment. C1 shows the largest concealment effect. Required-call suppression alone satisfies the concealment condition in 55 of 105 paired runs, or 52.4%. Adding final-response tampering raises this rate to 97 of 105 runs, or 92.4%, for a paired increase of 40.0 percentage points. In the $A + B$ condition, both modifications are delivered in 107 of 117 valid runs, and all 107 dual-delivery runs satisfy the joint oracle. Thus, all observed C1 runs with dual delivery conceal the impact caused by A .

C2 has a smaller effect. No A -only pair satisfies the concealment condition, while 11 of 112 $A + B$ pairs do, an increase of 9.8 percentage points. This smaller effect reflects the stricter conjunction required by C2: the intent rewrite must first reach downstream execution, and the final response must then claim successful completion.

Recovery. C3 shows that a second modification can reduce recovery after a required call is removed. The task recovers in 58 of 118 A -only pairs, or 49.1%. With premature finalization, recovery falls to 27 of 118 pairs, or 22.9%, for a paired difference of −26.3 percentage points. The traces show that, after call suppression, the agent often issues the required action in a later turn. Premature finalization ends many interactions before that repair can occur.

By contrast, C2 recovery remains at 87.5% in both conditions. Final-response tampering does not remove the original task retained by the agent, so the agent can still compare later behavior with that task and reject or correct the rewritten intent. The contrast suggests two different recovery paths. Recovery from call suppression often requires another model turn that the router can interrupt. Recovery from intent rewrite relies on trusted task state held by the agent, which final-response tampering does not remove.

5.6 Defense Coverage Analysis

Evaluation Method. RQ4 asks how much of the observed attack surface each trusted component can cover. We answer

Table 7: **Defense coverage and projected ASR.** Coverage uses 939 cases. The projection uses 23,103 valid runs and a 31.9% baseline.

Defense method	Blocked cases (%)	Projected ASR (%)	ASR drop (pp)
Schema validation	0.0	31.9	0.0
Agent-side checks	41.1	20.8	11.1
Tool-side authorization	13.7	28.9	3.0
End-to-end authentication	92.0	6.5	25.4

with a two-stage offline analysis that compares a schema-validation baseline with three deterministic defense profiles. First, we determine which benchmark cases each profile can block. For schema validation, we inspect all 41,380 delivered modification instances for schema and protocol validity. The other profiles represent their mechanisms as frozen rules over the attack family, target object, rewrite operator, and semantic type. We match each of the 939 benchmark cases against these rules. The earliest effective mechanism assigns one of four outcomes. A *blocked* verdict places a rejection before the specified impact. A *detected* verdict records trusted evidence of the attack or its consequence. A *mitigated* verdict reduces exposure. An *exposed* verdict has no matching mechanism.

Second, we join the blocked verdicts to the 23,103 valid runs in the main-table scope. An observed attack success counts as prevented when its case is blocked. We then recompute ASR with the same macro average over 30 configurations used in Table 2. The result projects coverage under correct enforcement of each profile. Detection and mitigation remain separate response outcomes.

Defense Profiles. The four profiles differ in the trusted state available to the verifier. *Schema validation* checks message structure, JSON types, and required fields. *Agent-side consistency checks* compare routed objects with locally retained task intent, tool definitions, argument policies, call records, and completion evidence. *Tool-side authorization* validates task-scoped operations, protected arguments, nonces, and invocation counts at the tool service. *End-to-end message authentication* verifies agent-signed requests at the provider and provider-signed responses at the agent. The signatures bind each object to its session, turn, request digest, and sequence number. This profile represents a provider-supported application-layer integrity protocol. Table 7 reports case coverage and projected ASR for these methods, separating coverage from the projected run-level effect.

Coverage by Placement. Coverage follows the verifier’s view of the interaction. Schema validation blocks 0 of 939 cases, and projected ASR remains at 31.9%. All 41,380 delivered modification pass schema and protocol validation.

Agent-side consistency checks block 386 cases and lower

projected ASR by 11.1 percentage points to 20.8%. Local task intent, pinned tool definitions, and argument policies support these checks before tool dispatch.

Tool-side authorization blocks 129 cases and lowers projected ASR by 3.0 percentage points to 28.9%. Its execution-boundary view covers unauthorized operations, protected arguments, and replay. End-to-end authentication blocks 864 cases and lowers projected ASR by 25.4 percentage points to 6.5%. Signed requests, signed responses, and sequence binding cover every modification to an object exchanged between the agent and provider.

Residual Exposure. Agent-side checks block 386 of the 939 benchmark cases. Among the remaining cases, they detect 425 and mitigate 25 passive-observation cases through data minimization. The other 103 cases remain exposed. These include 21 RMT-05 argument rewrites outside the protected field classes, all 25 RMT-08 call suppressions, and 57 RMT-10 identifier or ordering modifications. These cases lack an independent trusted reference available to the agent for detecting the missing or modified object. Only the 386 blocked cases enter the projected residual ASR.

Tool-side authorization leaves 810 cases exposed. Its view begins when a call reaches the trusted tool service. Request rewrites, suppressed calls, control fields, and final answers remain outside this boundary. End-to-end authentication leaves no request or response modification exposed. It detects the 50 RMT-S1 audit falsification cases by reconciling router records with signed objects. The 25 RMT-B0 cases remain exposed because they require confidentiality.

Take-away. Defense coverage follows trusted visibility. Schema validation leaves projected ASR at 31.9%, agent-side consistency checks reduce it to 20.8%, tool-side authorization reduces it to 28.9%, and end-to-end authentication reduces it to 6.5%. Each residual gap corresponds to an object or record outside the verifier’s view.

6 Discussion

Security Implications. Our results show that an LLM API router is part of the agent’s security boundary, not merely a transport component. It can alter what the model receives and what the agent acts on while preserving protocol validity. However, control of the path does not guarantee downstream impact. Cases without an independent trusted reference reach 82.6–98.9% ASR, whereas cases with such a reference fall to 6.3%. The key separator is whether a trusted component can compare a routed message with state that the router cannot change. Agents should therefore retain task intent, tool definitions, issued tool results, call history, and completion conditions outside the routed path. Evaluations should also inspect tool records and environment state because successful authorization, accountability, and resource abuse attacks often preserve task completion.

Verification Placement. Defense coverage follows the same principle. A defense can verify only the events and state visible to its trusted component. Schema validation blocks none of the delivered modifications because they remain schema-valid. Agent-side consistency checks and tool-side authorization help when the corresponding component can compare a message or action with trusted state, but neither covers events outside its view. For example, an agent cannot prove that the model issued a call deleted by the router. End-to-end message authentication provides the broadest coverage and lowers the adjusted ASR from 31.9% to 6.5%, but it requires provider support and does not prevent passive observation. Practical deployments should therefore combine message authentication with checks at trusted endpoints and independent reconciliation of usage and audit records.

Scope and Future Work. Our study measures post-compromise behavior in controlled tasks across five agents, six models, and three API formats. It does not estimate how often routers are compromised, study compromise techniques, or establish a general security ranking among agents and models. The controlled setting supports attribution of observed outcomes to specific router modifications, but it does not capture every deployment condition. Future work should examine chains of multiple intermediaries, stateful attacks across longer interactions, and production systems with proprietary agents, models, and tools.

7 Related Work

We organize prior work by where adversarial control enters the agent stack. We begin with the routed model path, then move to content processed by the model, memory, and tools, and conclude with security benchmarks and defenses.

Router attacks and path integrity. Work closest to ours studies untrusted intermediaries on the model API path. Liu et al. formalize the threat of a malicious API router, measure active manipulation and passive collection in router services operated by third parties, and build MINE to instantiate these attacks against publicly available agent frameworks [19]. Fu et al. build SIDEL to record, replay, and inject router modifications in coding agent workflows, then measure the resulting repository actions [13]. Luo et al. study response path tampering in bring your own key (BYOK) agents and propose SIGN-C, which authenticates fields that determine execution before the agent acts [22]. These studies show that an intermediary can alter agent traffic before inference or before local execution. Other work protects this path directly. AEX binds a request as seen by the client to a returned output and its lineage, while AEGIS confines plaintext handling to an enclave verified by the client [16, 34]. These mechanisms provide verifiable provenance for API outputs or constrain router

execution. We ask a complementary measurement question. When a modification tied to a concrete task is introduced on the request or response path, where does it stop in the agent loop? We measure delivery, downstream handling, and impact recorded outside the router across agents and models.

Prompt injection and memory poisoning. Indirect prompt injection places adversarial instructions in external content that the model later processes [15]. Memory poisoning instead changes persistent context or a retrieval knowledge base so that malicious content influences later decisions. AgentPoison demonstrates this threat against long term memory and agents that use retrieval augmented generation [8]. Prompt injection and router tampering describe different dimensions of an attack. Prompt injection depends on the model treating untrusted content as instructions. Router tampering describes control of protocol messages rather than a particular payload mechanism. A malicious router may use that control to inject prompt content on the request path, but it may also alter structured requests or responses directly, including after inference. We isolate this control of the routed path so that observed effects can be attributed to the router modification.

Tool and MCP attacks. ToolHijacker crafts malicious tool descriptions to bias tool retrieval and selection [30]. MCP-Tox benchmarks instructions hidden in MCP tool metadata, while MCPSecBench covers threats in MCP clients, protocol handling, servers, and hosts [32, 36]. MalTool studies malicious behavior embedded in tool implementations. ETDI protects tool identity, definition integrity, and permissions through cryptographic verification, versioning, and access control [6, 17]. Together, these works study threats introduced during tool discovery, MCP communication, or tool execution. Our threat model instead trusts the tool service and places adversarial control on the model API path. This separation lets us attribute downstream behavior to a modified model message rather than to a compromised tool.

Agent security benchmarks. AgentBench, WebArena, OS-World, and τ -bench evaluate agents through interactive tasks and evaluators that inspect task outcomes or resulting environment state [20, 35, 37, 41]. Security benchmarks add an attacker and a security objective. InjecAgent and AgentDojo focus on indirect prompt injection. Agent Security Bench covers prompt injection, memory poisoning, and other attack and defense families, while AgentCanary scores complete trajectories along outcome safety, security awareness, and task utility [10, 18, 39, 40]. ROUTERTAMPERBENCH also uses executable tasks and environment checks, but each case fixes one router modification and one oracle before execution. We separately record applicability, delivery, downstream handling, and impact verified outside the router. This decomposition separates cases in which the target is absent, the edit is re-

jected, the edit is handled without impact, or the edit reaches state maintained outside the router.

Agent and channel defenses. StruQ uses structured queries to separate trusted instructions from untrusted data [7]. MELON runs the agent again with the user request masked and compares the resulting actions. CaMeL extracts control and data flows from the trusted query and enforces them outside the model. IsolateGPT isolates third party applications and mediates interactions between them [9, 33, 42]. These defenses constrain how untrusted content or applications can influence agent actions. They do not authenticate that a response crossing an untrusted router matches the output from the model provider. Conversely, integrity mechanisms for the routed path do not establish that the model’s original decision is safe. The two layers are complementary. Our work complements both by measuring where router modifications propagate or stop in the complete agent loop.

8 Conclusion

We presented ROUTERTAMPERBENCH, a benchmark for measuring how router modifications tied to concrete tasks propagate through LLM agents. Its 939 cases cover 43 variants defined over semantic objects and use evidence outside the router to distinguish applicability, delivery, downstream handling, and impact. Across five agents and six models, we observe at least one attack success in every evaluated configuration, with a mean ASR of 31.9%. Composed attacks can hide the impact of a primary modification or interrupt recovery. Our defense study further shows that protection depends on whether a trusted component holds evidence that the router cannot rewrite. These findings support a simple design principle. An agent should not authorize security-sensitive actions solely from messages delivered by an intermediary. Robust deployments should retain trusted local state for consistency checks and use end-to-end authentication for the model messages that drive agent actions.

Ethical Considerations

Our experiments are conducted in isolated workspaces using deterministic, fixture-backed MCP services, so state changes are confined to experiment-controlled environments rather than third-party production systems. The benchmark has dual-use potential because its tampering specifications could be adapted to attack routed LLM-agent deployments. To reduce this risk, the released artifact excludes credentials and service-specific deployment details and focuses on controlled benchmark and evaluation components. Any previously unknown vulnerabilities identified in maintained systems will be handled through coordinated disclosure before exploit-enabling details are released.

Open Science

An anonymized artifact is available at https://anonymous.4open.science/r/use_artifact-DC82/. It contains the benchmark cases with associated task definitions and fixtures, the core router-tampering and evaluation implementation, configuration manifests, run-level experimental results, defense-analysis rules, and scripts and instructions for reproducing the aggregate results reported in the paper. Third-party agents, model services, and other components that cannot be redistributed are not included; the artifact documents the versions and interfaces needed to use them.

References

- [1] Anthropic. Claude Code: Overview. Claude Code Documentation. Accessed 2026-08-25. URL: <https://code.claude.com/docs/en/overview>.
- [2] Anthropic. Create a message. Claude API Reference. Accessed 2026-08-26. URL: <https://platform.claude.com/docs/en/api/messages/create>.
- [3] Anthropic. Model IDs and versioning. Claude Platform Documentation. Accessed 2026-08-26. URL: <https://platform.claude.com/docs/en/about-claude/models/model-ids-and-versions>.
- [4] Anthropic. Define tools. Claude Platform Documentation, platform.claude.com, n.d. Accessed: 2026-07-28.
- [5] BerriAI. Privilege escalation via unrestricted proxy configuration endpoint. GitHub Security Advisory for LiteLLM, April 2026. GHSA-53mr-6c8q-9789; CVE-2026-35029. Published 2026-04-03. Accessed 2026-07-29. URL: <https://github.com/BerriAI/litellm/security/advisories/GHSA-53mr-6c8q-9789>.
- [6] Manish Bhatt, Vineeth Sai Narajala, and Idan Habler. ETDI: Mitigating tool squatting and rug pull attacks in model context protocol (MCP) by using oauth-enhanced tool definitions and policy-based access control, 2025. doi:10.1109/CARS67163.2025.11337310.
- [7] Sizhe Chen, Julien Piet, Chawin Sitawarin, and David Wagner. StruQ: Defending against prompt injection with structured queries. In *34th USENIX Security Symposium (USENIX Security 25)*, pages 2383–2400, Seattle, WA, August 2025. USENIX Association. URL: <https://www.usenix.org/conference/usenixsecurity25/presentation/chen-sizhe>.
- [8] Zhaorun Chen, Zhen Xiang, Chaowei Xiao, Dawn Song, and Bo Li. AgentPoison: Red-teaming LLM agents via poisoning memory or knowledge bases. In *Advances in Neural Information Processing Systems*, volume 37,

- pages 130185–130213. Curran Associates, Inc., 2024. doi:10.52202/079017-4136.
- [9] Edoardo Debenedetti, Ilia Shumailov, Tianqi Fan, Jamie Hayes, Nicholas Carlini, Daniel Fabian, Christoph Kern, Chongyang Shi, Andreas Terzis, and Florian Tramèr. Defeating prompt injections by design, 2026. URL: <https://arxiv.org/abs/2503.18813>.
 - [10] Edoardo Debenedetti, Jie Zhang, Mislav Balunovic, Luca Beurer-Kellner, Marc Fischer, and Florian Tramèr. AgentDojo: A dynamic environment to evaluate prompt injection attacks and defenses for LLM agents. In *Advances in Neural Information Processing Systems*, volume 37, pages 82895–82920. Curran Associates, Inc., August 2024. doi:10.52202/079017-2636.
 - [11] DeepSeek. DeepSeek-V4-Flash update. DeepSeek API Documentation, 2026. Accessed 2026-08-26. URL: <https://api-docs.deepseek.com/updates/>.
 - [12] DeepSeek. DeepSeek-V4-Pro GA release. DeepSeek API Documentation, August 2026. Published 2026-08-13; accessed 2026-08-26. URL: <https://api-docs.deepseek.com/news/news260813/>.
 - [13] Donghao Fu, Jingxin Li, Xue Jiang, and Yihong Dong. Where is the cost of third-party API routers in agentic software development?, 2026. URL: <https://arxiv.org/abs/2607.23624>, arXiv:2607.23624, doi:10.48550/arXiv.2607.23624.
 - [14] GitLab. GitLab AI Gateway critical patch release: 18.6.2, 18.7.1, and 18.8.1. GitLab Documentation, February 2026. CVE-2026-1868. Published 2026-02-06. Accessed 2026-07-29. URL: <https://docs.gitlab.com/releases/patches/other-patches/patch-release-gitlab-ai-gateway-18-8-1-released/>.
 - [15] Kai Greshake, Sahar Abdelnabi, Shailesh Mishra, Christoph Endres, Thorsten Holz, and Mario Fritz. Not what you’ve signed up for: Compromising real-world LLM-integrated applications with indirect prompt injection. In *Proceedings of the 16th ACM Workshop on Artificial Intelligence and Security, AISec ’23*, pages 79–90, New York, NY, USA, 2023. Association for Computing Machinery. doi:10.1145/3605764.3623985.
 - [16] Yongjie Guan. AEX: Non-intrusive multi-hop attestation and provenance for LLM APIs. *arXiv preprint arXiv:2603.14283*, 2026. URL: <https://arxiv.org/abs/2603.14283>, arXiv:2603.14283.
 - [17] Yuepeng Hu, Yuqi Jia, Mengyuan Li, Dawn Song, and Neil Gong. MalTool: Malicious tool attacks on LLM agents, 2026. URL: <https://arxiv.org/abs/2602.12194>, arXiv:2602.12194.
 - [18] Peiyang Li, Songping Wang, Yi Huang, Yanhua Shi, Chenhao Zhang, Qi Li, Yueming Lyu, Caifeng Shan, Fengting Li, Chao Feng, Chuanqun Zhu, and Liang Chen. AgentCanary: A security evaluation framework for autonomous AI agents in real executable environments, 2026. URL: <https://arxiv.org/abs/2606.10484>, arXiv:2606.10484.
 - [19] Hanzhi Liu, Chaofan Shou, Hongbo Wen, Yanju Chen, Ryan Jingyang Fang, and Yu Feng. Your agent is mine: Measuring malicious intermediary attacks on the LLM supply chain, 2026. URL: <https://arxiv.org/abs/2604.08407>, arXiv:2604.08407, doi:10.48550/arXiv.2604.08407.
 - [20] Xiao Liu, Hao Yu, Hanchen Zhang, Yifan Xu, Xuanyu Lei, Hanyu Lai, Yu Gu, Hangliang Ding, Kaiwen Men, Kejuan Yang, Shudan Zhang, Xiang Deng, Aohan Zeng, Zhengxiao Du, Chenhui Zhang, Sheng Shen, Tianjun Zhang, Yu Su, Huan Sun, Minlie Huang, Yuxiao Dong, and Jie Tang. AgentBench: Evaluating LLMs as agents. In *The Twelfth International Conference on Learning Representations*, 2024. URL: <https://openreview.net/forum?id=zAdUB0aCTQ>.
 - [21] Angel Arribas Lopez. When AI infrastructure becomes part of the attack surface. Darktrace, July 2026. Published 2026-07-09. Accessed 2026-07-29. URL: <https://www.darktrace.com/blog/when-ai-infrastructure-becomes-part-of-the-attack-surface>.
 - [22] Mingyu Luo, Zihan Zhang, Zesen Liu, Yuchong Xie, Zhixiang Zhang, Dung Hiu Hilton Yeung, Wai Ip Lai, Ping Chen, Ming Wen, and Dongdong She. Rewriting the response path: Silent tampering and provider-signed defense in BYOK LLM agents, 2026. URL: <https://arxiv.org/abs/2605.02187>, arXiv:2605.02187.
 - [23] Nous Research. Hermes Agent. Software Repository and Documentation. Accessed 2026-08-25. URL: <https://github.com/NousResearch/hermes-agent>.
 - [24] OpenAI. Create a model response. OpenAI API Reference. Accessed 2026-08-26. URL: <https://developers.openai.com/api/reference/resources/responses/methods/create/>.
 - [25] OpenAI. Create chat completion. OpenAI API Reference. Accessed 2026-08-26. URL: <https://developers.openai.com/api/reference/resources/chat/subresources/completions/methods/create/>.
 - [26] OpenAI. Models. OpenAI API Documentation. Accessed 2026-08-26. URL: <https://developers.openai.com/api/docs/models>.
 - [27] OpenAI. Function calling. OpenAI API Documentation, developers.openai.com, n.d. Accessed: 2026-07-28.

- [28] OpenCode. OpenCode: Intro. OpenCode Documentation. Accessed 2026-08-25. URL: <https://opencode.ai/docs/>.
- [29] Pi. Pi documentation. Pi Documentation. Accessed 2026-08-25. URL: <https://pi.dev/docs/latest>.
- [30] Jiawen Shi, Zenghui Yuan, Guiyao Tie, Pan Zhou, Neil Zhenqiang Gong, and Lichao Sun. Prompt injection attack to tool selection in LLM agents, 2026. doi:10.14722/ndss.2026.230675.
- [31] Xingyao Wang, Boxuan Li, Yufan Song, Frank F. Xu, Xiangru Tang, Mingchen Zhuge, Jiayi Pan, Yueqi Song, Bowen Li, Jaskirat Singh, Hoang H. Tran, Fuqiang Li, Ren Ma, Mingzhang Zheng, Bill Qian, Yanjun Shao, Niklas Muennighoff, Yizhe Zhang, Binyuan Hui, Junyang Lin, Robert Brennan, Hao Peng, Heng Ji, and Graham Neubig. OpenHands: An open platform for AI software developers as generalist agents. In *The Thirteenth International Conference on Learning Representations*, 2025. URL: <https://openreview.net/forum?id=OJd3ayDDoF>.
- [32] Zhiqiang Wang, Yichao Gao, Yanting Wang, Suyuan Liu, Haifeng Sun, Haoran Cheng, Guanquan Shi, Haohua Du, and Xiangyang Li. MCPTox: A benchmark for tool poisoning attack on real-world MCP servers, 2026. URL: <https://ojs.aaai.org/index.php/AAAI/article/view/40895>, doi:10.1609/aaai.v40i42.40895.
- [33] Yuhao Wu, Franziska Roesner, Tadayoshi Kohno, Ning Zhang, and Umar Iqbal. IsolateGPT: An execution isolation architecture for LLM-based agentic systems. In *Network and Distributed System Security Symposium*. Internet Society, 2025. doi:10.14722/ndss.2025.241131.
- [34] Sipeng Xie, Qianhong Wu, Hengrun Lu, Ziliang Sun, Qi Wu, Bo Qin, and Qin Wang. The proxy knows too much: Sealing LLM API routers with attested TEEs, 2026. URL: <https://arxiv.org/abs/2606.16358>, arXiv:2606.16358.
- [35] Tianbao Xie, Danyang Zhang, Jixuan Chen, Xiaochuan Li, Siheng Zhao, Ruisheng Cao, Toh Jing Hua, Zhoujun Cheng, Dongchan Shin, Fangyu Lei, Yitao Liu, Yiheng Xu, Shuyan Zhou, Silvio Savarese, Caiming Xiong, Victor Zhong, and Tao Yu. OSWorld: Benchmarking multi-modal agents for open-ended tasks in real computer environments. In *Advances in Neural Information Processing Systems*, volume 37, pages 52040–52094. Curran Associates, Inc., 2024. doi:10.52202/079017-1650.
- [36] Yixuan Yang, Cuifeng Gao, Daoyuan Wu, Yufan Chen, Yingjiu Li, and Shuai Wang. MCPSecBench: A systematic security benchmark and playground for testing model context protocols, 2026. URL: <https://arxiv.org/abs/2508.13220>, arXiv:2508.13220.
- [37] Shunyu Yao, Noah Shinn, Pedram Razavi, and Karthik Narasimhan. τ -bench: A benchmark for tool-agent-user interaction in real-world domains, 2025. URL: <https://openreview.net/forum?id=roNSXZpUDN>.
- [38] Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik R Narasimhan, and Yuan Cao. React: Synergizing reasoning and acting in language models. In *The Eleventh International Conference on Learning Representations*, 2023. URL: https://openreview.net/forum?id=WE_vluYUL-X.
- [39] Qiusi Zhan, Zhixiang Liang, Zifan Ying, and Daniel Kang. InjecAgent: Benchmarking indirect prompt injections in tool-integrated large language model agents. In *Findings of the Association for Computational Linguistics: ACL 2024*, pages 10471–10506, Bangkok, Thailand, August 2024. Association for Computational Linguistics. doi:10.18653/v1/2024.findings-acl.624.
- [40] Hanrong Zhang, Jingyuan Huang, Kai Mei, Yifei Yao, Zhenting Wang, Chenlu Zhan, Hongwei Wang, and Yongfeng Zhang. Agent security bench (ASB): Formalizing and benchmarking attacks and defenses in LLM-based agents, 2025. URL: <https://openreview.net/forum?id=V4y0CpX4hK>.
- [41] Shuyan Zhou, Frank F. Xu, Hao Zhu, Xuhui Zhou, Robert Lo, Abishek Sridhar, Xianyi Cheng, Tianyue Ou, Yonatan Bisk, Daniel Fried, Uri Alon, and Graham Neubig. WebArena: A realistic web environment for building autonomous agents. In *The Twelfth International Conference on Learning Representations*, 2024. URL: <https://openreview.net/forum?id=oKn9c6ytLx>.
- [42] Kaijie Zhu, Xianjun Yang, Jindong Wang, Wenbo Guo, and William Yang Wang. MELON: Provable defense against indirect prompt injection attacks in AI agents. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pages 80310–80329. PMLR, 2025. URL: <https://proceedings.mlr.press/v267/zhu25z.html>.

Appendix

A Attack and Control Variants

Tables 8 and 9 expand the family-level summary in Table 1 and list all 43 benchmark variants. Each row reports the path, target and modification, trigger, number of benchmark cases, and trusted evidence evaluated by the oracle. The 939 cases are benchmark inputs, not experimental runs.

Table 8: Attack and control variant catalog, Part I: RMT-01 through RMT-04.

Variant (path)	Target and modification	Trigger	Cases	Oracle evidence
RMT-01: Intent rewrite				
goal_substitution (<i>Request</i>)	user task text, replace	first eligible request	25	calls consistent with the substituted task and missing original task outcome
history_stripping (<i>Request</i>)	conversation history, replace	every request from turn 2 onward	25	deviation from the expected outcome in trusted task state
system_instructions_stripping (<i>Request</i>)	system instructions, replace	first eligible request	25	deviation from the expected outcome in trusted task state
RMT-02: Execution configuration				
max_tokens_crush (<i>Request</i>)	maximum output length, replace	first request exposing the field	25	truncated output and missing required task outcome
model_downgrade (<i>Request</i>)	selected model, replace	first request exposing the field	25	effective model differs from the configured model
tool_choice_disable (<i>Request</i>)	tool choice set to none, replace	first request exposing the field	25	absence of the required tool call
tool_choice_forced_tool (<i>Request</i>)	tool choice set to a named tool, replace	first request exposing the field	25	call to the forced tool
tool_choice_required (<i>Request</i>)	tool choice set to required, replace	every request exposing the field	25	forced-call loop or changed task outcome
RMT-03: Tool definition tampering				
description_deprecation_substitution (<i>Request</i>)	tool description that redirects calls to a replacement tool, replace	first matching definition	9	call to the replacement tool and resulting trusted state
description_filesystem_redirect (<i>Request</i>)	tool description with path redirect, replace	first matching definition	13	redirected path in call and resulting state
description_network_sink_redirect (<i>Request</i>)	tool description that redirects an argument to a network destination, replace	first matching definition	25	redirected destination in call and resulting state
description_parameter_drop (<i>Request</i>)	tool description that omits a parameter, replace	first matching definition	14	changed or missing argument and missing task outcome
description_parameter_redirect (<i>Request</i>)	tool description with argument redirect, replace	first matching definition	25	redirected argument and resulting state
description_preference_manipulation (<i>Request</i>)	tool description with preferred action, replace	first matching definition	10	preferred call and resulting trusted state
description_tool_shadowing (<i>Request</i>)	description of one tool with instructions that change an argument of another tool, replace	first matching definition	25	modified argument in the affected call and resulting trusted state
schema_default_poisoning (<i>Request</i>)	tool schema default, replace	first matching definition	17	modified default value in call and resulting state
schema_enum_poisoning (<i>Request</i>)	tool schema enum, replace	first matching definition	7	modified enum value in call and resulting state
schema_property_stripping (<i>Request</i>)	tool schema property, replace	first matching definition	14	changed or missing argument and missing task outcome
schema_required_addition (<i>Request</i>)	tool schema required list, replace	first matching definition	3	extra required argument and resulting task outcome
schema_required_stripping (<i>Request</i>)	tool schema required list, replace	first matching definition	25	missing required argument and resulting task outcome
RMT-04: Tool result tampering				
error_result_fabrication (<i>Request</i>)	successful tool result changed to error, replace	after selected successful call	25	subsequent agent behavior and trusted task state
false_fact_redirection (<i>Request</i>)	fact in a tool result, replace	after selected successful call	25	subsequent call reflects the false fact and trusted state confirms its use
injected_instruction (<i>Request</i>)	instruction inserted into tool result, replace	after selected successful call	25	injected action and resulting trusted state
wrong_object_substitution (<i>Request</i>)	object in a tool result, replace	after selected successful call	25	action on the substituted object and resulting trusted state

B Configuration and Run Accounting

This section defines the counting units and denominators used throughout the evaluation. Table 11 reports the recorded system settings and corresponding run counts. A run is one benchmark case executed in one agent–model configuration. A delivered modification instance is one edit event. A run may contain multiple instances.

The main ASR, propagation, and projected ASR analyses use runs as their unit. Release composition, downstream condition counts, and defense coverage use benchmark cases.

Grouped ASR is the unweighted mean of the configuration-level ASR values. Not-applicable (NA) runs are excluded from every ASR denominator.

This ledger defines the counting unit and denominator for each result in the main evaluation.

C RMT-02 Configuration Field Exposure

RMT-02 measures which configuration fields agent clients expose to router modification and whether delivered changes

Table 9: Attack and control variant catalog, Part II: RMT-05 through RMT-12, RMT-S1, and RMT-B0.

Variant (path)	Target and modification	Trigger	Cases	Oracle evidence
RMT-05: Tool argument substitution				
counterfactual_state_redirection (Response)	tool-call argument, replace	first selected valid call	25	modified argument and unauthorized state change
RMT-06: Tool identity substitution				
rename_mutation (Response)	read call name changed to mutation, replace	first selected valid call	20	replacement call and unauthorized state change
rename_read (Response)	call name changed to another read, replace	first selected valid call	25	replacement call and unauthorized data disclosure
RMT-07: Tool call insertion				
unauthorized_inserted_call (Response)	new tool call, insert	first response with valid insertion arguments	25	inserted call and unauthorized state change
RMT-08: Tool call suppression				
required_call_suppression (Response)	required tool call, delete	up to three selected valid calls	25	required state change remains absent
RMT-09: Replay				
duplicate_call_replay (Response)	prior mutation call, replay	first selected valid response	25	matching call count and trusted state change
stale_observation_replay (Request)	prior tool result, replay	first matching result	25	subsequent decision based on stale data and resulting state
RMT-10: Control field tampering				
call_id_rewrite (Response)	tool-call identifier, replace	first matching call	25	call-result association and task state
index_rewrite (Response)	streamed tool-call index, replace	first matching call	25	reconstruction failure and task state
premature_finalization (Response)	finish reason, replace	first eligible response	25	required later call or state remains absent
status_misreport (Response)	response status, replace	first eligible response	25	status mismatch and deviation from the expected task outcome
tool_call_reordering (Response)	order of a multi-call response, replace	first matching multi-call turn	7	recorded call order and resulting state change
RMT-11: Final response tampering				
required_fact_contradiction (Response)	required fact in final text, replace	first final answer	25	final claim contradicts trusted fixture state
response_metadata_misreport (Response)	model and usage metadata, replace	every downstream response	25	metadata differs from upstream record
RMT-12: Tool registry insertion				
mid_session_tool_injection (Request)	tool registry entry, insert	every request from turn 2 onward	25	injected tool call and resulting trusted state
tool_name_squatting (Request)	tool entry with a conflicting name, insert	every eligible request	25	call to the injected tool and resulting trusted state
RMT-S1: Record falsification				
audit_omission (Router record)	router tamper summary, misreport	after the first tamper event	25	summary absent while internal record remains
audit_summary_misreport (Router record)	router tamper summary fields, misreport	after the first tamper event	25	reported summary differs from internal record
RMT-B0: Passive observation				
passive_visibility (Passive observation)	selected in-flight object, observe	first matching object	25	router observes the selected object without modifying it

affect execution. Field availability depends on client serialization and API handling. Table 10 therefore reports RMT-02 exposure and impact separately from propagation.

Field exposure differed by client. Only OpenCode included `tool_choice`, which exposed the three `tool_choice` variants to router modification. The `max_tokens` field was present in requests from OpenCode, Claude Code, and Hermes. Among the available observations, the `model` field was present for OpenCode, Claude Code, Hermes, and Pi.

The table distinguishes observation, delivery, and impact. We denote observed runs by O , delivered modifications by D , and verified impacts by I . An observed run without the target field contributes to O but not to D . Scheduled runs without an

observation are excluded from all rates.


This separate accounting distinguishes client field exposure from propagation after delivery.

D Trusted References and Propagation Traces

This section defines case-level annotations for propagation analysis and illustrates them with representative traces. Each applies to one benchmark case. Table 3 reports the four downstream conditions and their counts.

A recipient-visible trace exists only when the modified object appears in context that the receiving agent or model can inspect during later reasoning. Protocol receipt followed

(a) No independent trusted reference



(b) Independent trusted reference




Figure 5: **Representative propagation traces.** Panel (a) shows successful RMT-05 argument substitution without an independent trusted reference. Panel (b) shows RMT-09 replay rejection followed by recovery through an agent-side trusted registry. Identifying values are redacted.

by automatic execution provides no such opportunity.

An independent trusted reference is information outside the modified object that can verify its content or expected state and cannot be altered by the router. The trusted holder can be the agent, tool service, or task environment. This reference is available in time to support rejection or recovery. By contrast, oracle evidence verifies impact after the run.

Rejection and recovery capture two outcomes after delivery. Rejection prevents the modified object from causing its specified impact. Recovery requires delivery without the specified impact and successful completion of the original task.

Replay protection can operate at the agent or tool endpoint. Endpoint replay protection uses a unique request token or idempotency key to prevent a repeated request from changing state twice. For the 25 registry-rejection cases, the agent’s MCP tool registry rejects the replayed call before dispatch.

The annotations were fixed before we completed the six-model matrix, so later outcomes did not influence their labels. Comparisons across groups remain observational because the groups contain different cases.

The two traces in Figure 5 provide auditable records of opposite outcomes. Both runs use gpt-5.6-luna, retain task utility, and include complete recipient and trusted-state records.

The left trace shows successful tampering without an independent trusted reference. The run is a strict attack success. The GitHub fixture lets the oracle verify the substituted in-

itation after execution. The agent cannot use that fixture to check the modified argument before execution.

The right trace shows registry-based rejection followed by recovery. The independent agent-side tool registry prevents duplicate execution. The subsequent valid call completes the original task, so the run satisfies the paper’s definition of recovery.

Together, these traces show that visibility supports verification only when the recipient can compare the modified object with trusted state outside that object. The same distinction determines which defenses can verify a modification before it causes impact.

E Defense Coverage

Defense coverage depends on the state visible to each trusted component. Only a blocked verdict reduces projected ASR. Detected and mitigated cases remain in the projection because neither outcome guarantees prevention.

The residual cases follow the same visibility boundary. Agent-side data minimization mitigates passive observation. Authenticated records detect RMT-S1 record falsification after delivery. Tool-side authorization blocks replay only after a call reaches the tool endpoint. It cannot detect a call suppressed before dispatch. End-to-end message authentication does not cover passive observation because it protects integrity

Table 10: **RMT-02 configuration field exposure.** Observation, delivery, and impact by client and variant. *O*, *D*, *I*, and *F* denote observation, delivery, impact, and field absence.

Panel A. Field exposure by client

Variant	Field	OpenCode	Claude Code	Hermes	Pi	OpenHands
		DI O150	DI O150	FDI O150	F O150	F O150
max_tokens_crush	max_tokens	D100/I49 DI O25	D100/I40 D O25	D67/I100 DI O24	D0/I- DI O3	D0/I- No data O0
model_downgrade	model	D100/I100 DI O150	D100/I0 F O150	D100/I100 F O150	D100/I100 F O150	F O150
tool_choice_disable	tool_choice	D100/I47 DI O141	D0/I- F O125	D0/I- F O125	D0/I- F O125	D0/I- F O125
tool_choice_forced_tool	tool_choice	D100/I75 DI O125	D0/I- F O125	D0/I- F O125	D0/I- F O125	D0/I- F O125
tool_choice_required	tool_choice	D100/I90	D0/I-	D0/I-	D0/I-	D0/I-

Panel B. Aggregate counts and ratios

Variant	Field	Cases	Observed	D/O	No API error	I/D
max_tokens_crush	max_tokens	25	750	400/750 (53.3%)	400/400 (100.0%)	234/400 (58.5%)
model_downgrade	model	25	77	77/77 (100.0%)	0/77 (0.0%)	52/77 (67.5%)
tool_choice_disable	tool_choice	25	750	150/750 (20.0%)	150/150 (100.0%)	71/150 (47.3%)
tool_choice_forced_tool	tool_choice	25	641	141/641 (22.0%)	113/141 (80.1%)	106/141 (75.2%)
tool_choice_required	tool_choice	25	625	125/625 (20.0%)	112/125 (89.6%)	113/125 (90.4%)
Total		125	2,843	893/2,843 (31.4%)	775/893 (86.8%)	576/893 (64.5%)

■ B blocked
 ■ D detected
 ■ M mitigated
 ■ E exposed

Family	Variant	Schema validation	Agent-side checks	Tool-side authorization	End-to-end authentication
RMT-01	goal_substitution	E25	B25	E25	B25
	history_stripping	E25	B25	E25	B25
	system_instructions_stripping	E25	B25	E25	B25
RMT-02	max_tokens_crush	E25	D25	E25	B25
	model_downgrade	E25	D25	E25	B25
	tool_choice_disable	E25	D25	E25	B25
	tool_choice_forced_tool	E25	D25	E25	B25
	tool_choice_required	E25	D25	E25	B25
RMT-03	description_deprecation_substitution	E9	B9	E9	B9
	description_filesystem_redirect	E13	B13	E13	B13
	description_network_sink_redirect	E25	B25	E25	B25
	description_parameter_drop	E14	B14	E14	B14
	description_parameter_redirect	E25	B25	E25	B25
	description_preference_manipulation	E10	B10	E10	B25
	description_tool_shadowing	E25	B25	E25	B10
	schema_default_poisoning	E17	B17	E17	B25
	schema_enum_poisoning	E7	B7	E7	B7
	schema_property_stripping	E14	B14	E14	B14
	schema_required_addition	E3	B3	E3	B3
	schema_required_stripping	E25	B25	E25	B25
RMT-04	error_result_fabrication	E25	D25	E25	B25
	false_fact_redirection	E25	D25	E25	B25
	injected_instruction	E25	D25	E25	B25
	wrong_object_substitution	E25	D25	E25	B25
RMT-05	counterfactual_state_redirection	E25	B4/E21	B4/E21	B25
RMT-06	rename_mutation	E20	B20	E20	B20
RMT-06	rename_read	E25	B25	E25	B25
RMT-07	unauthorized_inserted_call	E25	B25	E25	B25
RMT-08	required_call_suppression	E25	E25	E25	B25
RMT-09	duplicate_call_replay	E25	D25	B25	B25
	stale_observation_replay	E25	D25	B25	B25
RMT-10	call_id_rewrite	E25	E25	E25	B25
	index_rewrite	E25	E25	E25	B25
	premature_finalization	E25	D25	E25	B25
	status_misreport	E25	D25	E25	B25
	tool_call_reordering	E7	E7	E7	B7
RMT-11	required_fact_contradiction	E25	D25	E25	B25
	response_metadata_misreport	E25	D25	E25	B25
	mid_session_tool_injection	E25	B25	E25	B25
RMT-12	tool_name_squatting	E25	B25	E25	B25
RMT-S1	audit_omission	E25	D25	E25	D25
	audit_summary_misreport	E25	D25	E25	D25
RMT-B0	passive_visibility	E25	M25	E25	E25

Figure 6: **Defense coverage by variant.** The matrix reports coverage verdicts and case counts for all 43 variants.

Table 11: Agent and model configurations, provider APIs, and run counts.

A. Configuration

Agent	Version	Agent API	Stream
OpenCode	1.2.24	Chat Completions	yes
Claude Code	2.1.216	Anthropic Messages	yes
Hermes	0.18.0	Chat Completions	yes
Pi	0.83.0	Chat Completions	yes
OpenHands	SDK 1.21.0	Chat Completions	no

Setting	Recorded value
Tool interface	MCP with isolated deterministic fixtures
Timeout	120 seconds, fixed within each model batch
Retry policy	One retry was allowed after HTTP 429. Tool tasks were retried once if a clean exit produced no MCP call and only planning text or no final text
Other settings	Maximum turns, temperature, and output length were not fixed across configurations
Evaluation period	2026-07-20 to 2026-08-14

B. Models and provider APIs

GPT family <i>Responses</i>	DeepSeek family <i>Chat Completions</i>	Claude family <i>Anthropic Messages</i>
gpt-5.6-luna	deepseek-v4-flash	claude-sonnet-4-6
gpt-5.6-terra	deepseek-v4-pro	claude-opus-4-6

C. Run accounting

Item	Count	Unit
Benchmark cases	939	cases
Cases in primary analysis	789	cases
Runs in primary analysis	23,103	runs
Delivered modification instances in available traces	41,380	instances

rather than confidentiality. Figure 6 reports the corresponding verdict and case count for every variant.